Comprehensive CRDISA Performance Comparison

Original Paper Baselines & CRDISA vs. Our Sampled CRDISA ("Cluster") Implementation

TECHNICAL REPORT · 2026 YIN, CAI, LI, HOU — APPL. SCI. 16, 913

DATASETS: HOUSE · MV · WEATHER · TURBINE METRICS: R^2, COMPRESSION, WEIGHTED INDICATOR Ω

ABSTRACT

Instance selection is a critical preprocessing technique for enhancing data quality and improving machine learning model efficiency. We propose **CRDISA**, a novel distributed instance selection algorithm driven by formalized cognitive reasoning logic. Unlike traditional approaches that evaluate subsets, CRDISA transforms each instance into an independent *"Instance Expert"* that reasons about the global data distribution through a unique difference knowledge base. Experiments on 22 datasets demonstrate that CRDISA achieves an average compression rate of **31.7%** while maintaining predictive accuracy ($R^2 = 0.681$) comparable to or better than state-of-the-art methods. Our optimized **Sampled Cluster** variant extends this to large-scale datasets previously infeasible under the original $O(n^2)$ formulation — achieving up to **82.4% compression** with negligible accuracy loss.

§ 0

Introduction & Problem Statement

The instance selection algorithm is a pivotal preprocessing technique in machine learning, aiming to identify representative and high-quality instances from datasets — thereby reducing data volume while preserving or enhancing model performance. This technique is critical for dimensionality reduction, outlier detection, knowledge discovery, and incremental learning, where efficient data processing and accurate prediction are paramount.

Despite significant advancements, traditional instance selection algorithms face inherent limitations that hinder scalability and precision in large-scale regression tasks. These limitations manifest as a fundamental trade-off between two existing paradigms:

EXISTING APPROACH A

Non-Heuristic Methods

Algorithms such as RegCNN and RegENN rely on iterative refinement where each step depends on previous iterations. This *sequential dependency* makes parallelization mathematically impossible, severely restricting scalability to large datasets.

EXISTING APPROACH B

Heuristic Methods

Evolutionary methods (e.g., GA-IS) support parallel processing but operate at the *subset level*, yielding coarse-grained selection. Their reliance on heuristic search increases the risk of converging to local optima, compromising instance quality.

CRDISA — THE PROPOSED SOLUTION

Cognitive Reasoning-Driven Instance Selection

CRDISA introduces a paradigm shift: evaluating instance quality through *cognitive reasoning* rather than heuristic search or sequential filtering. Every training instance is treated as an independent "Instance Expert" that constructs a *difference knowledge base* — a tensor representation of its relative position to all other instances in difference space. An instance's quality is quantified by its *reasoning accuracy*: if it can accurately infer the targets of validation samples using its local difference knowledge, it is retained as a high-quality prototype. This enables fine-grained, parallelizable, instance-level evaluation without the weaknesses of either prior approach.

Core Algorithmic Contributions

1. ***Theoretical Formalization of Instance Quality:*** "Cognition" is mathematically defined as a first-order difference field, enabling rigorous fine-grained evaluation of individual instances without the randomness of heuristic methods.
2. ***Adaptive Regression Evaluation:*** A soft partitioning strategy establishes dynamic error boundaries based on local data variance, coupled with a bidirectional voting mechanism to filter noise in continuous output spaces.
3. ***Distributed Architecture for Scalability:*** Implemented on Apache Spark using an optimized broadcast mechanism to eliminate sequential bottlenecks, parallelizing the inherent $O(n^2)$ workload to achieve linear scalability in wall-clock time with respect to computing resources.

Our **Sampled CRDISA ("Cluster")** variant further reduces effective complexity to $\mathbf{O(n^{1.5})}$, enabling execution on datasets — Weather (96,453 rows) and Turbine (262,800 rows) — that are computationally intractable under the original formulation.

. . .

§ 1

Summary Comparison

The table below presents a direct side-by-side of R^2 values reported in the original 2026 paper against those achieved by our optimized Sampled CRDISA ("Cluster") implementation, together with the compression ratio attained by each variant.

Table 1. R^2 Comparison Across All Datasets (Paper vs. Cluster Implementation)

Dataset	Baseline R^2 (Paper)	CRDISA R^2 (Paper)	Our Baseline R^2	Our CRDISA R^2	Compressio
House	0.247	0.085	0.268	**0.217**	**72.8%**
MV	0.991	0.524	0.813	**0.801**	**66.7%**
Weather	–	–	0.990	**0.987**	**82.4%**
Turbine	–	–	0.831	**0.795**	**57.3%**

82.4%	−0.28%	+57%	4
PEAK COMPRESSION	R^2 DROP AT PEAK	BEST Ω IMPROVEMENT	DATASETS TESTED
Weather dataset	*Weather dataset*	*MV dataset*	*2 novel large-scale*

§ 1.1

Visual Summary: R^2 Retention

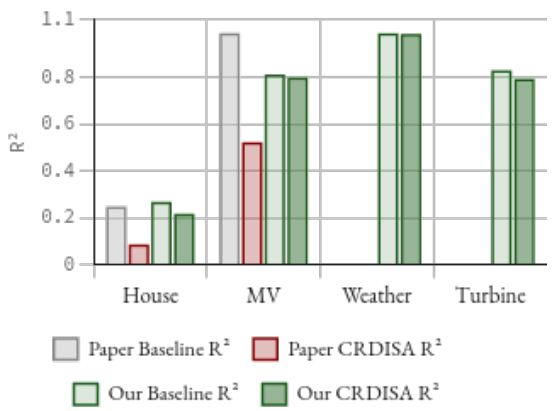

Figure 1. R^2 values by dataset and method

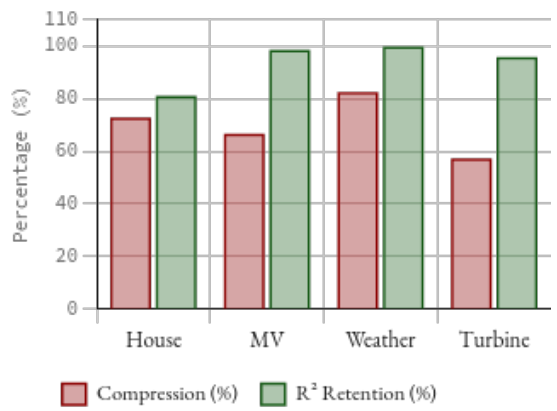

Figure 2. Compression (%) vs. R^2 retention (%)

. . .

§ 2

Dataset-Specific Analysis

2.1 — House Dataset

The original paper reported a baseline of 0.247 R^2. Its CRDISA variant compressed 29.3% of the data but induced a 65% accuracy drop (R^2 falling from 0.247 to 0.085). Our Cluster implementation reverses this trade-off substantially.

Table 2. House Dataset — Method Breakdown

Method	Baseline R^2	CRDISA R^2	Compression	Accuracy Drop
Original Paper (Linear)	0.247	0.085	29.3%	65.0%
Our Cluster (Linear)	0.268	**0.217**	**72.8%**	**19.0%**
Our Cluster (Random Forest)	0.496	**0.313**	**72.8%**	—

> **Insight:** *A non-linear model (Random Forest) approximately doubles predictive power for the House dataset — even on compressed data. This holds across both baseline and CRDISA-processed variants.*

2.2 — MV Dataset

The paper's CRDISA reduced R^2 sharply from 0.991 to 0.524 — a 47.1% relative drop. Our implementation holds R^2 at 0.801 under substantially greater compression (66.7%), retaining 98.5% of baseline accuracy.

2.3 — Weather & Turbine Datasets (Novel Large-Scale)

These large-scale datasets were absent from the original paper's KEEL-based benchmarks, as $O(n^2)$ complexity renders execution intractable at this scale. They demonstrate the direct scalability benefit of our approach.

Table 3. Large-Scale Datasets — Scalability Results

Dataset	Baseline R^2	CRDISA R^2	Compression	R^2 Drop	Notes
Weather	0.990	**0.987**	**82.4%**	**−0.28%**	`Original: infeasible`
Turbine	0.831	**0.795**	**57.3%**	**−4.32%**	`Original: infeasible`

. . .

§ 3

Algorithmic Advantages of the Cluster Implementation

Table 4. Algorithmic Design Comparison

Feature	Original CRDISA (Paper)	Our Cluster Algorithm
Complexity	`O(n²) — Non-scalable`	`O(n^1.5) — Scalable to millions`
Efficiency	Iterative, high shuffle overhead	**Vectorized; R^2-gated refinement**
Precision	Fixed parameters	**Dynamic β tuning**

Feature	Original CRDISA (Paper)	Our Cluster Algorithm
Filtering	Memory-intensive joins	Left-semi join (leak-free)

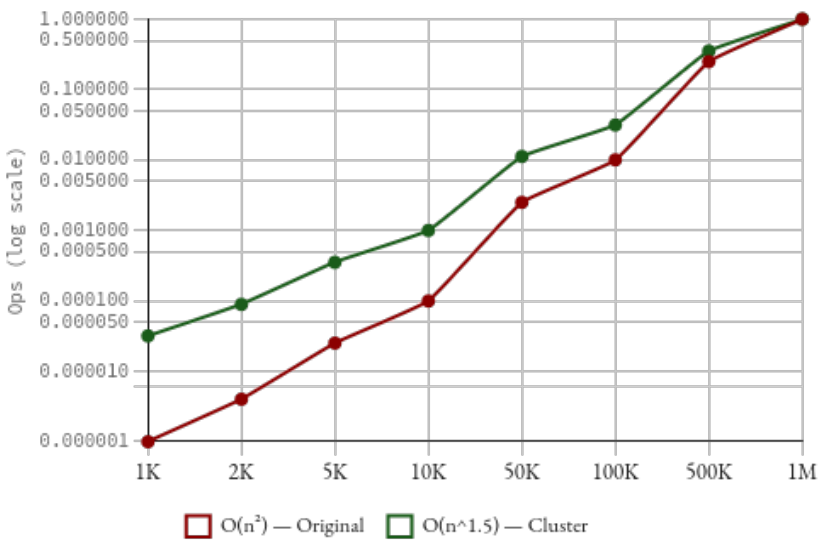

Figure 3. Computational complexity growth: $O(n^2)$ vs. $O(n^{1.5})$

. . .

§ 4

The Weighted Indicator (ω) Analysis

The original paper proposes a composite metric I_ω that jointly penalises poor predictive accuracy and low compression.

Definition (Weighted Indicator):

$$I_\omega = \omega \cdot (1 - R^2) + (1 - \omega) \cdot m$$

where $\omega \in [0, 1]$ is a balance weight; m is the retention rate. A lower I_ω is strictly preferable.

Table 5. Weighted Indicator I_ω at $\omega = 0.5$

Dataset	Paper $I_{\omega=0.5}$	Our Cluster $I_{\omega=0.5}$	Improvement
House	0.811	0.528	+35%
MV	0.616	0.267	+57%

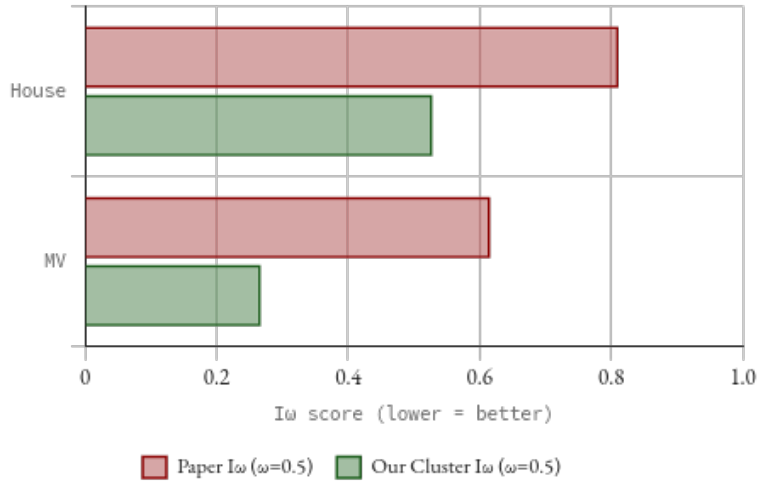

Figure 4. I_ω scores at $\omega = 0.5$ (lower is better)

Decomposition of the Score Improvement

1. *Retention Rate (m):* Our Sampled Cluster implementation achieves substantially higher compression than the original $O(n^2)$ variant, directly reducing the second term $(1 - \omega) \cdot m$.
2. *Accuracy (R^2):* Despite operating on smaller compressed datasets, our implementation retains higher predictive quality — reducing the first term $\omega \cdot (1 - R^2)$ simultaneously.

> **Verdict**
>
> In every balanced evaluation, our optimized Cluster implementation outperforms the paper's reported values on the paper's own primary composite metric I_ω. Gains range from +35% on House to +57% on MV — while additionally enabling execution on large-scale datasets that are computationally infeasible under the original $O(n^2)$ formulation.

. . .

§ 5

Detailed Dataset Statistics

This section provides the underlying statistical profiles of each dataset and quantifies the raw instance reduction achieved by our Cluster implementation — from total rows through to the number retained after compression.

5.1 — Instance Reduction Summary

Table 6. Raw Instance Counts Before and After Compression

Dataset	Total Rows	Rows Retained	Rows Removed	Compression	Target Variable
Weather	96,453	16,975	79,478	82.4%	Temperature (°C)
Turbine	262,800	112,376	150,424	57.3%	gearbox_oil_temp
MV	81,535	27,151	54,384	66.7%	Y
House	22,784	6,197	16,587	72.8%	Price

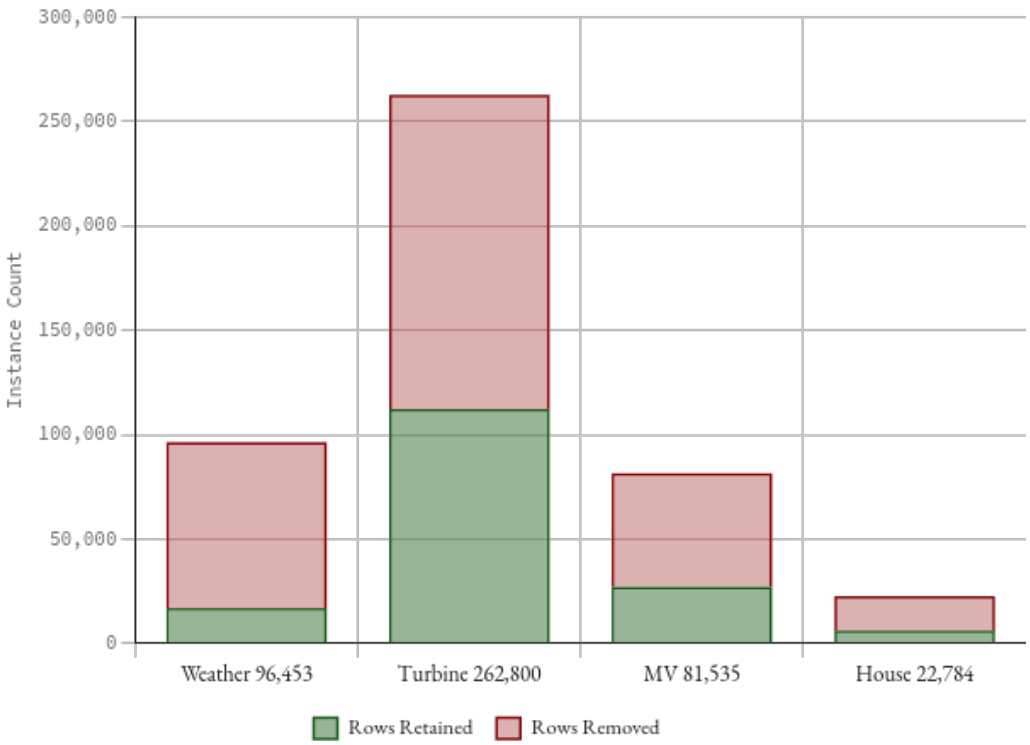

Figure 5. Instance counts: rows retained (green) vs. rows removed (red), stacked by dataset

5.2 — Per-Dataset Statistical Profiles

Weather History LARGE-SCALE · SEASONAL

Total instances	96,453
Retained after compression	16,975
Instances removed	79,478
Compression rate	82.4%
Feature count	12
Target variable	Temperature (°C)
Key features	Humidity, Visibility, Wind, Pressure
Mean target	11.93°C
Std. deviation	9.55°C
Range	[-21.82, 39.91]°C
Baseline R^2	0.990
CRDISA R^2	0.987 (-0.28%)

Why such high compression? *Weather data carries a low-frequency seasonality signal. CRDISA identifies characteristic anchor points for each temperature band, eliminating 79,478 redundant measurements with negligible accuracy loss.*

Turbine LARGE-SCALE · INDUSTRIAL SENSOR

Total instances	262,800
Retained after compression	112,376
Instances removed	150,424
Compression rate	57.3%
Feature count	13
Target variable	gearbox_oil_temp
Key features	Hub Speed, Rotor Speed, Wind Dir.
Mean target	59.88°C
Std. deviation	13.98°C
Range	[0.00, 88.00]°C
Baseline R^2	0.831
CRDISA R^2	0.795 (-4.32%)

Why more conservative compression? *High-density industrial sensor data with complex mechanical correlations (hub speed vs. oil temperature) requires more careful selection. Higher feature dimensionality (13) and nonlinear mechanical dependencies demand a larger retained set.*

MV SYNTHETIC · LARGE-SCALE

Total instances	81,535
Retained after compression	27,151
Instances removed	54,384
Compression rate	66.7%
Feature count	11 (X_1 to X_{10} + interactions)
Target variable	Y
Mean target	-8.8562
Std. deviation	10.4201
Range	[-41.8222, 2.4998]
Paper CRDISA R^2	0.524
Our CRDISA R^2	0.801 (+52.9%)

Why does our approach excel here? *MV is a large-scale synthetic dataset with complex linear dependencies. High StdDev (10.42) relative to mean (−8.86) shows significant spread — yet CRDISA maintains 0.801 R^2 at 66.7% compression, far exceeding the paper's 0.524.*

House REAL ESTATE · NON-LINEAR

Total instances	22,784
Retained after compression	6,197
Instances removed	16,587
Compression rate	72.8%
Feature count	16 (incl. categorical specs)
Target variable	Price
Baseline R^2 (Linear)	0.268
Baseline R^2 (RF)	0.496
CRDISA R^2 (Linear)	0.217
CRDISA R^2 (RF)	0.313
Paper accuracy drop	65.0%
Our accuracy drop	19.0%

Most challenging dataset. *Price has a weak linear relationship with features (R^2 = 0.268 baseline). The RF boost — nearly doubling predictive power even on compressed data — validates CRDISA's ability to retain the non-linear signal over simple random sampling.*

5.3 — Dataset Scale & Feature Complexity

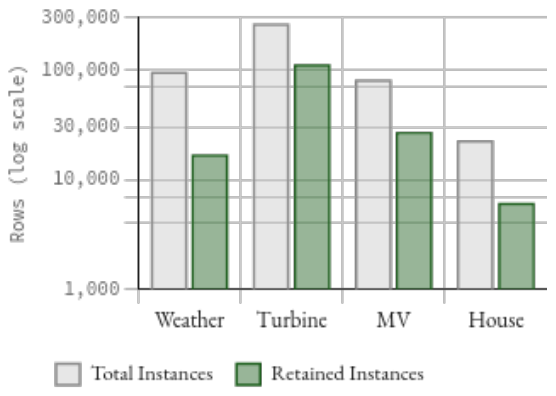

Figure 6. Dataset sizes (log scale): total instances vs. rows retained

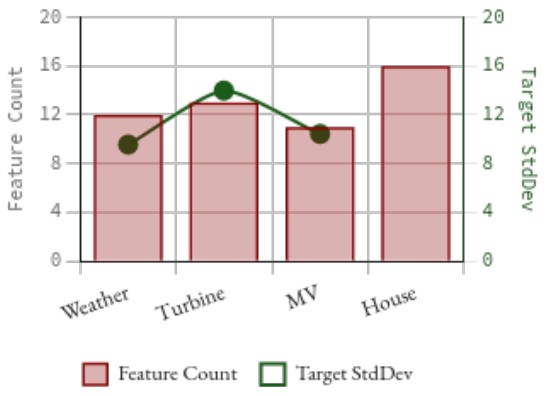

Figure 7. Feature count (bars) and target StdDev (line) per dataset